

# Fraud\_Detection

March 11, 2022

## 1 Credit Card Fraud Detection

It is important that credit card companies are able to recognize fraudulent credit card transactions so that customers are not charged for items that they did not purchase.

The dataset contains transactions made by credit cards in September 2013 by European cardholders. This dataset presents transactions that occurred in two days, where we have 492 frauds out of 284,807 transactions. The dataset is highly unbalanced, the positive class (frauds) account for 0.172% of all transactions.

It contains only numerical input variables which are the result of a PCA transformation. Unfortunately, due to confidentiality issues, we cannot provide the original features and more background information about the data. Features V1, V2, ... V28 are the principal components obtained with PCA, the only features which have not been transformed with PCA are 'Time' and 'Amount'. Feature 'Time' contains the seconds elapsed between each transaction and the first transaction in the dataset. The feature 'Amount' is the transaction Amount. Feature 'Class' is the response variable and it takes value 1 in case of fraud and 0 otherwise.

We will choose a model based on the precision and recall curve (and the area under it), which are more insightful metrics than accuracy in such an un-balanced dataset.

```
[1]: import pandas as pd

df=pd.read_csv('creditcard.csv')

df.shape
```

```
[1]: (284807, 31)
```

## 2 Splitting the data into training and test sets:

```
[2]: from sklearn.model_selection import train_test_split

df_train, df_test = train_test_split(df, test_size=0.2, random_state=42,
    ↳stratify=df['Class'])
```

### 3 Exploring the data:

```
[3]: df_train.head(3)
```

```
[3]:
```

	Time	V1	V2	V3	V4	V5	V6	\
265518	161919.0	1.946747	-0.752526	-1.355130	-0.661630	1.502822	4.024933	
180305	124477.0	2.035149	-0.048880	-3.058693	0.247945	2.943487	3.298697	
42664	41191.0	-0.991920	0.603193	0.711976	-0.992425	-0.825838	1.956261	

	V7	V8	V9	...	V21	V22	V23	\
265518	-1.479661	1.139880	1.406819	...	0.076197	0.297537	0.307915	
180305	-0.002192	0.674782	0.045826	...	0.038628	0.228197	0.035542	
42664	-2.212603	-5.037523	0.000772	...	-2.798352	0.109526	-0.436530	

	V24	V25	V26	V27	V28	Amount	Class
265518	0.690980	-0.350316	-0.388907	0.077641	-0.032248	7.32	0
180305	0.707090	0.512885	-0.471198	0.002520	-0.069002	2.99	0
42664	-0.932803	0.826684	0.913773	0.038049	0.185340	175.10	0

```
[3 rows x 31 columns]
```

```
[4]: df_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 227845 entries, 265518 to 17677
Data columns (total 31 columns):
#   Column  Non-Null Count  Dtype
---  -
0   Time    227845 non-null  float64
1   V1       227845 non-null  float64
2   V2       227845 non-null  float64
3   V3       227845 non-null  float64
4   V4       227845 non-null  float64
5   V5       227845 non-null  float64
6   V6       227845 non-null  float64
7   V7       227845 non-null  float64
8   V8       227845 non-null  float64
9   V9       227845 non-null  float64
10  V10      227845 non-null  float64
11  V11      227845 non-null  float64
12  V12      227845 non-null  float64
13  V13      227845 non-null  float64
14  V14      227845 non-null  float64
15  V15      227845 non-null  float64
16  V16      227845 non-null  float64
17  V17      227845 non-null  float64
18  V18      227845 non-null  float64
19  V19      227845 non-null  float64
```

```

20 V20      227845 non-null float64
21 V21      227845 non-null float64
22 V22      227845 non-null float64
23 V23      227845 non-null float64
24 V24      227845 non-null float64
25 V25      227845 non-null float64
26 V26      227845 non-null float64
27 V27      227845 non-null float64
28 V28      227845 non-null float64
29 Amount   227845 non-null float64
30 Class    227845 non-null int64
dtypes: float64(30), int64(1)
memory usage: 55.6 MB

```

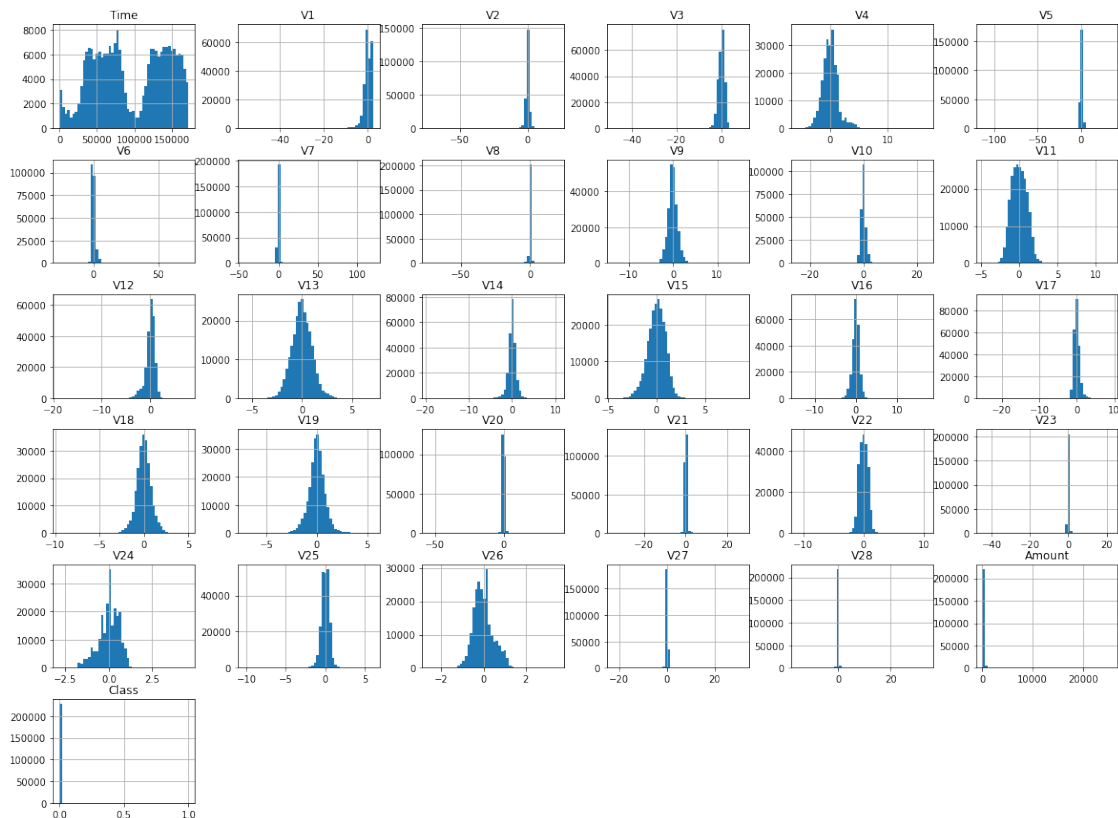
### 3.1 Plotting features' histograms:

```

[5]: import matplotlib.pyplot as plt

df_train.hist(bins=50, figsize=(20,15))
plt.show()

```



### 3.2 Calculating correlations between the labels and the features:

```
[6]: corr_matrix = df_train.corr()  
     corr_matrix["Class"].sort_values(ascending=False)
```

```
[6]: Class      1.000000  
     V11      0.153709  
     V4       0.135014  
     V2       0.090586  
     V21      0.035588  
     V19      0.032380  
     V8       0.020552  
     V20      0.019385  
     V27      0.016034  
     V28      0.009810  
     Amount   0.006211  
     V26      0.004119  
     V22      0.002926  
     V25      0.001618  
     V23     -0.004169  
     V15     -0.005705  
     V13     -0.005861  
     V24     -0.007483  
     Time     -0.010564  
     V6      -0.043334  
     V5      -0.093578  
     V9      -0.098247  
     V1      -0.100041  
     V18     -0.108732  
     V7      -0.186184  
     V16     -0.193826  
     V3      -0.194135  
     V10     -0.217894  
     V12     -0.259989  
     V14     -0.301054  
     V17     -0.321937  
     Name: Class, dtype: float64
```

## 4 Data preparations:

```
[7]: import numpy as np  
  
     X_train=df_train.to_numpy()[ : , :30]  
     y_train=df_train['Class'].to_numpy()  
  
     X_test=df_test.to_numpy()[ : , :30]  
     y_test=df_test['Class'].to_numpy()
```

```
[8]: def display_scores(scores):
      print("Scores:", scores)
      print("Mean:", scores.mean())
      print("Standard deviation:", scores.std())
```

## 5 Training an XGBoost model:

We will apply early stopping with the ‘aucpr’ metric which is the area under the precision-recall curve, and when fine-tuning we will use the ‘average\_precision’ scoring which is the same.

Setting aside a validation set:

```
[9]: X_train_xgb, X_val_xgb, y_train_xgb, y_val_xgb = train_test_split(X_train,
      ↪y_train, test_size=0.2,
      ↪random_state=42, stratify=y_train)
```

### 5.1 Fine-tuning:

```
[10]: from sklearn.model_selection import GridSearchCV
      import xgboost as xgb

      xgb_clf = xgb.XGBClassifier(n_estimators=100, objective='binary:logitraw',
      ↪random_state=42, use_label_encoder=False)

      fit_params={'verbose': False,
                  'eval_metric': 'aucpr',
                  'eval_set': [(X_val_xgb, y_val_xgb)],
                  'early_stopping_rounds': 10}

      param_grid = [{'max_depth': [15, 20, 25],
                     'learning_rate': [0.2, 0.5, 0.7]}]

      grid_search = GridSearchCV(xgb_clf, param_grid, cv=5,
      ↪scoring='average_precision', return_train_score=True, n_jobs=-1)
      grid_search.fit(X_train_xgb, y_train_xgb, **fit_params)
```

```
[10]: GridSearchCV(cv=5,
                  estimator=XGBClassifier(base_score=None, booster=None,
                                          colsample_bylevel=None,
                                          colsample_bynode=None,
                                          colsample_bytree=None, gamma=None,
                                          gpu_id=None, importance_type='gain',
                                          interaction_constraints=None,
                                          learning_rate=None, max_delta_step=None,
                                          max_depth=None, min_child_weight=None,
                                          missing=nan, monotone_constraints=None,
```

```

n_estimators=100, n_jobs=None,
num_parallel_tree=None,
objective='binary:logitraw',
random_state=42, reg_alpha=None,
reg_lambda=None, scale_pos_weight=None,
subsample=None, tree_method=None,
use_label_encoder=False,
validate_parameters=None, verbosity=None),
n_jobs=-1,
param_grid=[{'learning_rate': [0.2, 0.5, 0.7],
              'max_depth': [15, 20, 25]}],
return_train_score=True, scoring='average_precision')

```

```
[11]: grid_search.best_params_
```

```
[11]: {'learning_rate': 0.2, 'max_depth': 15}
```

## 5.2 Plotting the Precision-Recall curve:

first we want 'clean' predictions on the training set, using `cross_val_predict`:

```
[12]: from sklearn.model_selection import cross_val_predict

xgb_clf_curve = xgb.XGBClassifier(max_depth=15, n_estimators=100,
    ↪ learning_rate=0.2, objective='binary:logitraw',
    random_state=42, use_label_encoder=False)

fit_params={'eval_metric': 'aucpr',
            'eval_set': [(X_val_xgb, y_val_xgb)],
            'early_stopping_rounds': 10}

y_scores = cross_val_predict(xgb_clf_curve, X_train_xgb, y_train_xgb, cv=5,
    ↪ method="predict_proba",
                                fit_params=fit_params, n_jobs=-1)

```

```
[13]: y_scores=y_scores[:,1]
```

With these scores we compute precision and recall for all possible thresholds:

```
[14]: from sklearn.metrics import precision_recall_curve

precisions, recalls, thresholds = precision_recall_curve(y_train_xgb, y_scores)

precisions.shape, recalls.shape, thresholds.shape

```

```
[14]: ((32863,), (32863,), (32862,))
```


```
[15]: precisions=precisions[:-1]
      recalls=recalls[:-1]

[16]: precisions.shape, recalls.shape, thresholds.shape

[16]: ((32862,), (32862,), (32862,))

[17]: def plot_precision_vs_recall(precisions, recalls):
      plt.plot(recalls, precisions, "b--")
      plt.xlabel('Recall')
      plt.ylabel('Precision')
      plt.title('Precision vs Recall', fontdict={'fontname': 'Comic Sans MS',
      ↪ 'fontsize': 20})

      plot_precision_vs_recall(precisions, recalls)
      plt.show()
```



### 5.3 Suggesting threshold alternatives:

#### 5.3.1 Choosing the primary elbow in the PR curve:

It makes the most sense, producing both high precision and high recall. Said elbow appears at recall value of 0.8. Hence:

```
[18]: np.where(recalls>=0.8)[0][-1]
```

```
[18]: 32591
```


The precision and recall we receive are:

```
[19]: precisions[32591], recalls[32591]
```

```
[19]: (0.9197080291970803, 0.8)
```

Looking at it on the PR curve - the red dot:

```
[20]: def plot_precision_vs_recall(precisions, recalls):  
    plt.plot(recalls, precisions, "b--")  
    plt.xlabel('Recall')  
    plt.ylabel('Precision')  
    plt.title('Precision vs Recall', fontdict={'fontname': 'Comic Sans MS',  
→ 'fontsize': 20})  
    plt.plot(recalls[32591], precisions[32591], 'r^--', label='1222')  
  
plot_precision_vs_recall(precisions, recalls)  
plt.show()
```



The corresponding threshold is:



```
[21]: thresholds[32591]
```

```
[21]: -0.47505212
```

### 5.3.2 Choosing the secondary elbow in the PR curve:

If precision is vital we will take the 'smaller' elbow at recall value of 0.69:

```
[22]: np.where(recalls>=0.69)[0][-1]
```

```
[22]: 32642
```

The precision and recall we receive are:


```
[23]: precisions[32642], recalls[32642]
```

```
[23]: (0.9819819819819819, 0.692063492063492)
```

Looking at it on the PR curve - the red dot:

```
[24]: def plot_precision_vs_recall(precisions, recalls):
        plt.plot(recalls, precisions, "b--")
        plt.xlabel('Recall')
        plt.ylabel('Precision')
        plt.title('Precision vs Recall', fontdict={'fontname': 'Comic Sans MS',
        ↪ 'fontsize': 20})
        plt.plot(recalls[32642], precisions[32642], 'r^--', label='1222')

plot_precision_vs_recall(precisions, recalls)
plt.show()
```



The corresponding threshold is:

```
[25]: thresholds[32642]
```

```
[25]: 2.6995952
```

For this project we go with option number 1, which is the primary elbow:

```
[26]: model_threshold=-0.47505212
```

## 5.4 Final model training:

```
[27]: xgb_clf = xgb.XGBClassifier(max_depth=15, learning_rate=0.2, n_estimators=100,
    ↪objective='binary:logitraw', random_state=42, use_label_encoder=False)

fit_params={'verbose': False,
            'eval_metric': 'aucpr',
            'eval_set': [(X_val_xgb, y_val_xgb)],
            'early_stopping_rounds': 10}

xgb_clf.fit(X_train_xgb, y_train_xgb, **fit_params)
```

```
[27]: XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
    colsample_bynode=1, colsample_bytrees=1, gamma=0, gpu_id=-1,
```

```

importance_type='gain', interaction_constraints='',
learning_rate=0.2, max_delta_step=0, max_depth=15,
min_child_weight=1, missing=nan, monotone_constraints='()',
n_estimators=100, n_jobs=12, num_parallel_tree=1,
objective='binary:logitraw', random_state=42, reg_alpha=0,
reg_lambda=1, scale_pos_weight=1, subsample=1,
tree_method='exact', use_label_encoder=False,
validate_parameters=1, verbosity=None)

```

The predictions will be done the following way:

```

[28]: def xgb_predict(X, model, model_threshold):
        return (model.predict_proba(X)[:,-1] >= model_threshold)

```

Computing Precision, Recall and F1 score on the the validation set:

```

[29]: from sklearn.metrics import recall_score, precision_score, f1_score

y_val_pred=xgb_predict(X_val_xgb, xgb_clf, model_threshold)

print("Precision:", precision_score(y_val_xgb, y_val_pred))
print("Recall:", recall_score(y_val_xgb, y_val_pred))
print("F1 Score:", f1_score(y_val_xgb, y_val_pred))

```

```

Precision: 0.921875
Recall: 0.7468354430379747
F1 Score: 0.8251748251748252

```

## 6 Testing the model on the test set:

Computing Precision, Recall and F1 score on the the test set:

```

[30]: y_test_pred=xgb_predict(X_test, xgb_clf, model_threshold)

print("Precision:", precision_score(y_test, y_test_pred))
print("Recall:", recall_score(y_test, y_test_pred))
print("F1 Score:", f1_score(y_test, y_test_pred))

```

```

Precision: 0.9032258064516129
Recall: 0.8571428571428571
F1 Score: 0.8795811518324608

```